
ComRank: Ranking Loss for Multi-Label Complementary Label Learning

Jing-Yi Zhu^{1,2}, Yi Gao^{1,2*}, Miao Xu³, Min-Ling Zhang^{1,2}

¹School of Computer Science and Engineering, Southeast University, Nanjing 210096, China

²Key Laboratory of Computer Network and Information Integration (Southeast University),
Ministry of Education, China

³The University of Queensland, Australia

{zhujingyi, gao_yi, zhangml}@seu.edu.cn,
{miao.xu}@uq.edu.au

Abstract

Multi-label complementary label learning (MLCLL) is a weakly supervised paradigm that addresses multi-label learning (MLL) tasks using complementary labels (i.e., irrelevant labels) instead of relevant labels. Existing methods typically adopt an unbiased risk estimator (URE) under the assumption that complementary labels follow a uniform distribution. However, this assumption fails in real-world scenarios due to instance-specific annotation biases, making URE-based methods ineffective under such conditions. Furthermore, existing methods underutilize label correlations inherent in MLL. To address these limitations, we propose **ComRank**, a ranking loss framework for MLCLL, which encourages complementary labels to be ranked lower than non-complementary ones, thereby modeling pairwise label relationships. Theoretically, our surrogate loss ensures Bayes consistency under both uniform and biased cases. Experiments demonstrate the effectiveness of our method in MLCLL tasks. The code is available at <https://github.com/JellyJamZhu/ComRank>.

1 Introduction

Multi-label learning (MLL) refers to a task where an instance is associated with multiple relevant labels, which has broad applications in real-world scenarios [Zhang and Zhou, 2014, Tang et al., 2023, Kou et al., 2024]. However, accurately labeling a large number of instances with all their true labels incurs high labor costs. To address this issue, weakly supervised learning for MLL has gained widespread attention in recent years, including *semi-supervised multi-label learning* [Liu et al., 2006, Niu et al., 2019], *multi-label learning with missing labels* [Sun et al., 2010, Wu et al., 2014], and *partial multi-label learning* (PML) [Xie and Huang, 2018, Zhang and Fang, 2020].

Multi-label complementary label learning (MLCLL) has recently emerged as a weakly supervised learning paradigm, which enables algorithms to learn from complementary labels instead of relevant labels to address the MLL problem. In MLCLL, each training instance is associated with a complementary label rather than relevant labels, where the complementary label specifies a label that the instance **does not** belong to. One application for MLCLL is privacy preservation, such as in sensitive surveys where respondents may hesitate to provide all truthful answers. By only asking them to exclude certain options, data collection becomes easier and more privacy-friendly, while also reducing labeling costs.

*Corresponding author

The conventional solution for MLCLL currently revolves around the unbiased risk estimator (URE), which is a powerful tool in weakly supervised learning that enables the accurate estimation of true classification risk. Specifically, these methods derive URE by assuming that complementary labels follow uniform distribution, where a URE can be constructed based on common MLL loss functions such as binary cross-entropy loss (BCE), mean squared error loss (MSE), and mean absolute error loss (MAE). With the above uniform assumption, Gao et al. [2023] firstly investigate the MLCLL problem and derive a URE. Moreover, they propose a gradient-friendly MLCLL loss to enhance gradient updating of the URE. However, this uniform assumption may fail in real-world scenarios, where annotators may provide complementary labels with biases influenced by the characteristics of the instances. Moreover, previous URE-based methods do not fully exploit label correlations, which are critical in MLL.

In MLCLL, based on the fact that complementary labels are known to be irrelevant, while non-complementary labels may include relevant ones, it is generally desirable for the predicted probabilities of complementary labels to rank lower than those of non-complementary labels. This intuition naturally aligns with the objective of ranking loss. Inspired by this observation, we propose a **complementary ranking loss** (called **ComRank**) framework for MLCLL, which encourages learning by enforcing this ranking constraint and capturing pairwise relationships between labels. Additionally, our framework uses an exponential loss as the surrogate loss, while the complementary ranking loss achieves Bayes consistency under both cases of uniform and biased complementary labels. This overcomes the limitation of the URE, which only has theoretical guarantees under uniform complementary labels. Furthermore, our proposed framework can directly capture label correlation information from the rankings, offering unique advantages in MLL. Outstanding experimental results demonstrate the effectiveness of our method. The main contributions of this paper are as follows:

- We theoretically analyze why existing URE-based methods cannot work well on complementary labels with a biased distribution. URE strongly depends on the uniform assumption of complementary labels, and fails to estimate the expected risk when the complementary label distribution shifts.
- We firstly introduce ranking loss into MLCLL and propose a complementary ranking loss framework, ComRank, for MLCLL. We theoretically prove that it possesses Bayes consistency under both uniform and biased complementary labels. Experiments on different complementary label distributions demonstrate the outstanding performance of our method.

The remainder of this paper is organized as follows: Section 2 summarizes the related work, and preliminaries are introduced in section 3. Discussion on URE with different complementary label distribution are provided in section 4. Section 5 and section 6 introduce ComRank and its experimental results. The last section 7 concludes our work.

2 Related Work

Multi-label learning. MLL is a classification task where each instance can be related to multiple labels simultaneously [Jia et al., 2023, Shi et al., 2024]. Considering label correlations during training, there has been extensive theoretical exploration of ranking in MLL. Gao and Zhou [2011] first defined the consistency of MLL and Li et al. [2017] introduced the concept of Bayes consistency into the context of ranking loss for MLL, proposing a surrogate loss proven to be consistent from the perspective of the Bayes prediction rule. Xie and Huang [2018] shows that in probabilistic MLL, ranking between positive and negative labels can help disambiguate false positives. Li et al. [2024] illustrates that missing labels can be assumed to rank between positive and negative labels in weakly supervised MLL settings. However, these studies on ranking loss are all focused on supervised scenarios and are not applicable to the MLCLL problem.

Complementary label learning (CLL). CLL was first proposed to solve the multi-class classification tasks, which aims to train a multi-class classifier using complementary labels. Ishida et al. [2017] initially derived a URE by modifying pairwise-comparison and one-versus-all losses for CLL under the assumption of uniform complementary label distribution. To get rid of trapping in specific losses, Ishida et al. [2019] extended a general URE framework that can accommodate arbitrary losses. Recognizing that the uniform assumption for generating complementary labels may fail to handle the real-world scenarios, various methods that diverge from uniform assumption have also been

investigated, such as Yu et al. [2018] relaxed CLL to biased complementary labels by estimating a transition matrix, and Gao and Zhang [2021] developed a discriminative model without the uniform assumption. The success of these methods is based on the multi-class scenarios, but they may not be suitable for the MLCLL problem.

Multi-label complementary label learning. MLCLL was first introduced by Gao et al. [2023] as a solution to the challenge of collecting and accurately annotating multi-label data. Under the uniform assumption (i.e., randomly sampled from irrelevant labels), Gao et al. [2023] derived a URE, which ensures the classifier learned through complementary labels converges to the optimal classifier in MLL. Furthermore, in later research, Gao et al. [2024] allowed the use of biased complementary labels to recover relevant labels through an estimated transition matrix, and Gao et al. [2025] further investigated the URE in the setting of multiple complementary labels. Unfortunately, these methods fail to consider the biased complementary label distribution, or do not provide theoretical guarantees to derive a URE under biased distribution.

Therefore, in this paper, we demonstrate why URE-based methods cannot remain unbiased across non-uniform complementary label distributions, making it impossible for these methods to accurately estimate the true classification risk in MLL. Furthermore, we propose a novel complementary ranking loss framework, ComRank, for MLCLL, which offers theoretical guarantees under both uniform and biased complementary label distributions.

3 Preliminaries

In MLL, $\mathcal{X} \subset \mathbb{R}^d$ and $\mathcal{Y} \subseteq \{0, 1\}^K$ represent the d -dimensional feature space and the label space with K labels, respectively. A multi-label sample can be represented as $(\mathbf{x}, Y) \in \mathcal{X} \times \mathcal{Y}$, where \mathbf{x} and Y follow the probability distribution $p(\mathbf{x}, Y)$. Let the training set be $D = \{(\mathbf{x}_i, Y_i) \mid 1 \leq i \leq N\}$. Here, Y can be written as a K -dimensional vector $\mathbf{y} = [y^1, y^2, \dots, y^K] \in \{0, 1\}^K$, where $y^k = 1$ indicates that the label k is related to \mathbf{x} . The objective of MLL is to train a multi-label classifier $\mathbf{g} : \mathcal{X} \rightarrow \mathbb{R}^K$ by minimizing the following expected risk $R(\mathbf{g})$, where $g_k \in [0, 1]$ is the predicted probability for the k -th element and \mathcal{L} denotes the common MLL loss:

$$\mathbf{g}^* = \arg \min_{\mathbf{g} \in G} R(\mathbf{g}) = \arg \min_{\mathbf{g} \in G} \mathbb{E}_{p(\mathbf{x}, Y)} [\mathcal{L}(\mathbf{g}(\mathbf{x}), Y)]. \quad (1)$$

In MLCLL, we define the complementary training set as $\bar{D} = \{(\mathbf{x}_i, \bar{y}_i) \mid 1 \leq i \leq N\}$, where $\bar{y}_i \in \{\mathcal{Y} - Y_i\}$. To ensure the validity of \bar{y}_i , Y_i cannot be \emptyset or \mathcal{Y} , so $|\mathcal{Y}| = 2^K - 2$. \mathbf{x} and $\bar{\mathbf{y}}$ follow a distribution $p(\mathbf{x}, \bar{\mathbf{y}})$. For convenience, $\bar{\mathbf{y}}$ can be represented as a K -dimensional vector $\bar{\mathbf{y}} = [\bar{y}^1, \bar{y}^2, \dots, \bar{y}^K] \in \{0, 1\}^K$, where $\bar{y}^k = 1$ indicates that the label k is the complementary label of \mathbf{x} . MLCLL aims to train a multi-label classification classifier $\bar{\mathbf{g}} : \mathcal{X} \rightarrow \mathbb{R}^K$, which can predict relevant labels for unseen instances. Let $\bar{\mathcal{L}}$ represent the MLCLL loss, and the optimal classifier $\bar{\mathbf{g}}^*$ is obtained by minimizing the expected risk of MLCLL:

$$\bar{\mathbf{g}}^* = \arg \min_{\bar{\mathbf{g}} \in \bar{G}} \bar{R}(\bar{\mathbf{g}}) = \arg \min_{\bar{\mathbf{g}} \in \bar{G}} \mathbb{E}_{p(\mathbf{x}, \bar{\mathbf{y}})} [\bar{\mathcal{L}}(\bar{\mathbf{g}}(\mathbf{x}), \bar{\mathbf{y}})]. \quad (2)$$

Before commencing the analysis, we first present the definitions of URE. URE is an important tool for weakly supervised method [Ishida et al., 2019, Feng et al., 2020], providing an accurate estimation of the true risk $R(\mathbf{g})$, i.e., $R(\mathbf{g}) = \bar{R}(\mathbf{g})$. Therefore, $\bar{R}(\mathbf{g})$ is referred as a URE of $R(\mathbf{g})$.

4 Complementary Label Distribution

Existing URE-based methods in MLCLL recover relevant labels and construct corresponding classifiers from complementary labels by making reasonable assumptions about the distribution of complementary labels. As a result, the complementary label distribution plays a crucial role in the modeling process. However, current assumption and theorem have certain limitations. Therefore, a systematic discussion and reasonable extension of them will be conducted in this section. We start from the uniform assumption and the existing URE derived from it.

Assumption 4.1. [Ishida et al., 2017, Gao et al., 2023] Uniform Distribution Assumption:

$$p(k = \bar{y} \mid \mathbf{x}_i) = \begin{cases} \frac{1}{K - |Y_i|}, & \text{if } k \notin Y_i, \\ 0, & \text{if } k \in Y_i. \end{cases} \quad (3)$$

Assumption 4.1 provides information that complementary labels are uniformly selected from the label space excluding relevant labels. This implies that each irrelevant label has an equal probability of being chosen as a complementary label by annotators. Based on this distribution, the contribution of $p(\mathbf{x}, Y)$ to $p(\mathbf{x}, \bar{y})$ is also uniform. As a result, a URE can be easily derived according to Assumption 4.1, which is shown as follows.

Theorem 4.2. [Gao et al., 2023] *URE under the uniform distribution: With $p(k = \bar{y}|\mathbf{x})$ defined in Assumption 4.1 and $R(\mathbf{g})$ defined in Eq. (1), the equality $\bar{R}(\mathbf{g}) = R_u(\mathbf{g})$ holds, where*

$$R_u(\mathbf{g}) = \mathbb{E}_{p(\mathbf{x}, \bar{y})} \left[\frac{1}{2^{K-1} - 1} \sum_{Y \subseteq \mathcal{Y}, \bar{y} \neq Y} \mathcal{L}(\mathbf{g}, Y) \right]. \quad (4)$$

However, in the real world, a uniform distribution may hardly cover actual scenarios. Some label combinations may be more likely to occur than others. Labels with a higher correlation to the relevant labels are closer to relevant labels and less likely to be chosen as complementary labels [Gao et al., 2024]. For example, when the relevant label is *water*, annotators are more likely to choose *desert* (a low co-occurrence label) rather than *lake* (a high co-occurrence label) as the complementary label. At the same time, *lake* is more likely to be closer to the relevant label compared to *desert*.

Therefore, we extend complementary labels from a uniform distribution to a biased distribution. Let matrix $\mathbf{L} = [l]_{|\mathcal{Y}| \times K}$ represents the correlation matrix, where \mathbf{L} has $|\mathcal{Y}| = 2^K - 2$ rows, with each row corresponding to a label set $Y \in \mathcal{Y}$. The element l_{Yk} represents the correlation between the label set Y and the k -th label. The closer Y is to the k -th label, the larger l_{Yk} is. Note that \mathbf{L} is used only for inference and does not appear in later computations.

Assumption 4.3. Biased Distribution Assumption:

$$p(k = \bar{y}|\mathbf{x}) = \begin{cases} \frac{z}{l_{Yk}}, & \text{if } k \notin Y, \\ 0, & \text{if } k \in Y, \end{cases} \quad (5)$$

where $z = \frac{1}{\sum_{k=1, k \notin Y}^{K-1} \frac{1}{l_{Yk}}}$, $l_{Yk} \propto p(k \in Y|\mathbf{x})$.

Assumption 4.3 summarizes the condition probabilities of a biased complementary label distribution. Here, z in Eq. (5) is the normalization factor, ensuring that $p(k = \bar{y}|\mathbf{x})$ forms a valid probability distribution. The correlation l_{Yk} in Eq. (5) is proportional to $p(k \in Y|\mathbf{x})$, meaning the correlation between label k and set Y is proportional to the conditional probability that label k belongs to set Y , given \mathbf{x} .

Besides, the design of previous work is generally based on Label-Dependent Assumption, that is: The complementary label \bar{y} is independent of the features \mathbf{x} conditioned on the relevant label set Y , i.e., $p(\bar{y}|Y) = p(\bar{y}|\mathbf{x}, Y)$ [Ishida et al., 2017, 2019, Gao et al., 2023]. This assumption does not adequately encompass real-world scenarios, as annotators subconsciously select labels that are not too similar based on the instance’s features, rather than the relevant labels, in the process of choosing complementary labels. Therefore, we adopt a more realistic Instance-dependent Assumption, which has been widely used in other weakly supervised scenarios [Xia et al., 2020, Chen et al., 2021, Kou et al., 2023].

Assumption 4.4. Instance-Dependent Assumption: Given an instance \mathbf{x} , the complementary label \bar{y} is independent of Y , i.e. $p(\bar{y}|\mathbf{x}, Y) = p(\bar{y}|\mathbf{x})$.

Assumption 4.4 is a fundamental premise regarding the relationship between complementary labels and instances, positing that the selection of a complementary label depends on the instance. In some scenarios—such as an image containing multiple animals—annotators may struggle to exclude all relevant categories based solely on features. However, our assumption is motivated by the observation that annotators often eliminate obviously irrelevant labels by inspecting the input instance (e.g., excluding "building" from an image showing animals), without needing to infer the full set of relevant labels. This behavior supports modeling \bar{y} as conditionally independent of Y given \mathbf{x} . We adopt this assumption as a tractable approximation that reflects limited annotator under biased complementary label distributions.

Subsequently, we investigate the URE under the Biased Distribution Assumption. To simplify the computation, we first express the probability distribution in Assumption 4.4 in matrix form as a bias

transition matrix, denoted by $\bar{\mathbf{L}} = [\bar{l}_{Yk}]_{|\mathcal{Y}| \times K} \in \mathbb{R}^{(2^K - 2) \times K}$. Here, $\bar{l}_{Yk} = p(k = \bar{y} | \mathbf{x})$ represents the probability of label k being selected as the complementary label when the relevant label set for \mathbf{x} is Y . Theorem 4.5 then provides the URE derived from the biased complementary labels.

Theorem 4.5. *URE under biased distribution: Given $p(k = \bar{y} | \mathbf{x})$ and $R(\mathbf{g})$, the equality $R(\mathbf{g}) = \bar{R}(\mathbf{g}) = R_b(\mathbf{g})$ holds when*

$$R_b(\mathbf{g}) = \mathbb{E}_{p(\mathbf{x}, \bar{y})} \left[\sum_{Y \subseteq \mathcal{Y}, \bar{y} \neq Y} \bar{l}_{Y\bar{y}}^+ \mathcal{L}(\mathbf{g}, Y) \right], \quad (6)$$

where $\bar{l}_{Y\bar{y}}^+$ belongs to the matrix $\bar{\mathbf{L}}^+ = [\bar{l}^+]_{|\mathcal{Y}| \times K}$, and $\bar{\mathbf{L}}^+$ is the Moore-Penrose pseudoinverse of $\bar{\mathbf{L}}^T$.

The proof is stated in Appendix B. Compared to the URE under uniform distribution, the construction of URE (Eq. (6)) under biased distribution is heavily dependent on the specific characteristics of the distribution. The uniform assumption ensures that each complementary label contributes equally to the distribution of the relevant labels. However, once there is a bias in the contributions, $p(\mathbf{x}, \bar{y})$ is no longer uniform, which inevitably causes the URE to change. Therefore, the URE derived under the uniform assumption can no longer accurately estimate the expected risk once the distribution of complementary labels changes. In other words, it becomes ineffective under the biased assumption.

5 The Proposed Framework

5.1 Complementary Ranking Loss

In addition to the lack of universality across different distributions, URE-based methods do not take label correlations into account, thus losing important information needed to solve the MLL problem. Since complementary labels are known to be irrelevant, while non-complementary labels may include relevant ones, enforcing lower scores for complementary labels may help distinguish likely-relevant labels. This intuitive motivation naturally aligns with the goal of ranking loss, making it a suitable choice for integration into MLCLL. By incorporating label correlations without significantly increasing computational complexity, ranking loss has proven to be an effective tool for capturing pairwise label correlations [Zhang and Fang, 2020, Zhang et al., 2018, Fürnkranz et al., 2008]. In MLL, the traditional ranking loss is defined as

$$\mathcal{L}(\mathbf{g}(\mathbf{x}), Y) = \sum_{k=1, k \in Y}^K \sum_{j=1, j \notin Y}^K \mathbb{I}[g_k(\mathbf{x}) < g_j(\mathbf{x})],$$

where $\mathbb{I}(\cdot)$ is the indicator function, which outputs 1 when the condition holds and otherwise 0. Inspired by the complementary 0-1 loss [Chou et al., 2020], we propose the **complementary ranking loss** for MLCLL:

$$\bar{\mathcal{L}}(\mathbf{g}(\mathbf{x}), \bar{y}) = \sum_{k=1, k \neq \bar{y}}^K \mathbb{I}[g_k(\mathbf{x}) < g_{\bar{y}}(\mathbf{x})].$$

Unlike URE-based methods, the complementary ranking loss does not rely on any assumption regarding complementary labels. Instead, complementary ranking loss directly compares the predicted probabilities between different labels. Similar to ranking loss in supervised MLL, it penalizes cases where the complementary label predicted probability is higher than that of a non-complementary label, because a complementary label is certainly not a relevant label, while a non-complementary label may be either relevant or irrelevant. Therefore, it is generally reasonable to assign lower scores to complementary labels, as they are less likely to be relevant compared to non-complementary labels. The rationale behind this design will be formally justified in the next subsection.

However, directly optimizing the complementary ranking loss is challenging, as it is typically NP-hard due to its non-convexity and discontinuity. Thus, a convex surrogate loss can be introduced to facilitate more efficient optimization, a common method in ranking loss methods:

$$\bar{\mathcal{L}}(\mathbf{g}(\mathbf{x}), \bar{y}) = \sum_{k=1, k \neq \bar{y}}^K \ell(g_{\bar{y}}(\mathbf{x}) - g_k(\mathbf{x})),$$

Algorithm 1 ComRank Algorithm

Input: \bar{D} : the complementary-label training set $\{(\mathbf{x}_i, \bar{y}_i)\}_{i=1}^n$ θ : the initial parameters of classifier g T : the number of epochs \mathcal{A} : an external stochastic optimization algorithm**Output:** g : learned multi-label classifier**Training Routine**

- 1: **for** $t = 1$ to T **do**
 - 2: Let \mathcal{L} be the risk,
 - 3: $\mathcal{L} = \frac{1}{n} \sum_{i=1}^n \{\bar{\mathcal{L}}_{\text{CR}}(g(\mathbf{x}_i), \bar{y}_i)\}$
 - 4: Set gradients $-\nabla_{\theta} \mathcal{L}$
 - 5: Update θ by \mathcal{A} with $-\nabla_{\theta} \mathcal{L}$
 - 6: **end for**
-

where ℓ represents the surrogate loss. In this paper, we propose a complementary ranking loss framework named **ComRank**, using an exponential function as the surrogate loss, shown as follows. The algorithm of applying ComRank can be referred from Algorithm 1.

$$\bar{\mathcal{L}}_{\text{CR}}(g(\mathbf{x}), \bar{y}) = \sum_{k=1, k \neq \bar{y}}^K \exp(g_{\bar{y}}(\mathbf{x}) - g_k(\mathbf{x})).$$

5.2 Bayes Consistency for ComRank

Bayes consistency is a desirable property of a loss function, ensuring that minimizing the expected loss leads to the Bayes prediction rule [Li et al., 2017, Cheng et al., 2010]. We say that a loss has Bayes consistency if it leads g to follow the Bayes prediction rule:

$$g_k^*(\mathbf{x}) = p(k \in Y | \mathbf{x}).$$

In contrast, URE only guarantees that the expected risk of MLCLL aligns with that of fully supervised MLL, without directly evaluating the classification results. Therefore, Bayes consistency is a more rigorous criterion than URE.

To verify the rationality of ComRank, it's necessary to demonstrate ComRank's theoretical soundness by establishing Bayes consistency. The analysis begins with the following lemma.

Lemma 5.1. *Under $\bar{\mathcal{L}}_{\text{CR}}(g(\mathbf{x}), \bar{y})$, $\bar{g}^k(\mathbf{x}) \geq \bar{g}^j(\mathbf{x})$ if and only if $p(k = \bar{y} | \mathbf{x}) \leq p(j = \bar{y} | \mathbf{x})$.*

The proof can be found in Appendix C. The conclusion of Lemma 5.1 is achieved through risk minimization, which is a fundamental result that applies to all distributional scenarios of MLCLL. Moreover, this reflects the relationship between the predictive probabilities given by $\bar{\mathcal{L}}_{\text{CR}}(g(\mathbf{x}), \bar{y})$ and the probability of becoming a complementary label. With Lemma 5.1, we can demonstrate that $\bar{\mathcal{L}}_{\text{CR}}(g(\mathbf{x}), \bar{y})$ exhibits Bayes consistency under both the uniform distribution (Assumption 4.1) and the biased distribution (Assumption 4.3).

Theorem 5.2. *Bayes Consistency for ComRank: For both uniform and biased complementary label distributions, $\bar{g}^k(\mathbf{x}) \geq \bar{g}^j(\mathbf{x})$ holds under $\bar{\mathcal{L}}_{\text{CR}}(g(\mathbf{x}), \bar{y})$ if and only if $p(k \in Y | \mathbf{x}) \geq p(j \in Y | \mathbf{x})$.*

The proof is stated in Appendix D. Theorem 5.2 shows that ComRank establishes a theoretical connection whereby the ranking between complementary and non-complementary labels can be transferred to the ranking between irrelevant and relevant labels. Next, we will provide experimental results to support its performance.

Table 1: *Average Precision* (mean \pm std) on the training data with uniform complementary labels. The best performance of each dataset is shown in **boldface**, where \downarrow / \uparrow indicates that smaller/larger values of metrics are better performance.

Methods	L-UW	CCMN	PMLMD	PARD	MAE	GDF	$R_u(g)$	ComRank
scene	.395 \pm .016	.458 \pm .019	.441 \pm .043	.740 \pm .020	.432 \pm .019	.759 \pm .012	.734 \pm .014	.780\pm.010
yeast	.685 \pm .018	.646 \pm .017	.695 \pm .023	.608 \pm .118	.698 \pm .018	.712 \pm .019	.679 \pm .016	.715\pm.019
enron	.375 \pm .037	.337 \pm .042	.537 \pm .092	.444 \pm .035	.427 \pm .047	.620 \pm .067	.411 \pm .038	.634\pm.068
rcv1-s1	.445 \pm .020	.409 \pm .028	.348 \pm .030	.468 \pm .089	.427 \pm .029	.471 \pm .057	.363 \pm .019	.491\pm.118
bibtex	.237 \pm .009	.259 \pm .025	.280 \pm .044	.447 \pm .020	.287 \pm .009	.614 \pm .017	.413 \pm .016	.658\pm.014
bookmark	.506 \pm .009	.397 \pm .028	.181 \pm .005	.549 \pm .010	.506 \pm .007	.619 \pm .006	.512 \pm .007	.628\pm.005
nuswideBoW	.451 \pm .010	.431 \pm .014	.457 \pm .025	.457 \pm .057	.466 \pm .011	.553 \pm .008	.595\pm.008	.585 \pm .010

Table 2: Summary of pairwise t-test for ComRank against other comparing approaches at 0.05 significance level on uniform datasets, showing in form of Win/Tie/Loss.

ComRank against	L-UW	CCMN	PMLMD	PARD	MAE	GDF	$R_u(g)$	in total
<i>One Error</i>	5/2/0	5/2/0	6/1/0	5/2/0	5/2/0	4/3/0	6/1/0	36/13/0
<i>Coverage</i>	6/1/0	7/0/0	5/2/0	6/1/0	6/1/0	4/3/0	6/1/0	40/9/0
<i>Ranking Loss</i>	6/1/0	7/0/0	6/1/0	6/1/0	6/1/0	4/3/0	6/1/0	41/8/0
<i>Average Precision</i>	6/1/0	6/1/0	6/1/0	6/1/0	5/2/0	4/3/0	6/0/1	39/9/1
in total	23/5/0	25/3/0	23/5/0	23/5/0	22/6/0	16/12/0	24/0/1	156/40/0

6 Experiments

6.1 Experimental Setup

Datasets. To fully verify the effectiveness of ComRank, we select seven multi-label datasets for experiments². The range of their data sizes is from 1702 to 269648, and data domains include text, biology, and images. Their details can be referred from Appendix A. We unify data preprocessing on these datasets. To comprehensively illustrate the experimental impact of the single complementary label, in line with previous studies [Gao et al., 2023, 2024, Hang and Zhang, 2024], for datasets with label space larger than 50, we extract the 15 most frequently occurring labels and delete instances that did not appear with these labels.

Data Processing. We use uniform and biased complementary labels to conduct experiments. Specifically, 1) *Uniform complementary labels*: Each instance x_i is equipped with a complementary label randomly selected from $\{\mathcal{Y} - Y_i\}$, where irrelevant labels have an equal probability of being chosen; 2) *Biased complementary labels*: The selection of the complementary label for x_i follows a biased rule: based on co-occurrence rates computed from the original dataset, labels with lower co-occurrence rates are more likely to be selected. The model is trained on data annotated with complementary labels, while the test data is labeled with relevant labels.

Baselines. ComRank is compared with seven recent competitive methods, including one CLL method: L-UW; one MLL method: CCMN; two PML methods: PMLMD and PARD; and three MLCLL methods: MAE, GDF and $R_u(g)$, which are shown in the following details:

- L-UW [Gao and Zhang, 2021]: A CLL method that incorporates a weighted loss based on empirical risk to enhance the prediction gap between potential relevant labels and complementary labels.
- CCMN [Xie and Huang, 2023]: A MLL method that leverages class-conditional multi-label noise for learning, constructing two unbiased estimators.
- PMLMD [Xie et al., 2021]: A PML method in the form of ranking loss, specially equipped with weight and meta-disambiguation to figure out candidate labels in partial multi-label learning label sets.

²Publicly available at <https://mulan.sourceforge.net/datasets-mlc.html>.

Table 3: *Average Precision* (mean \pm std) on the training data with biased complementary labels. The best performance of each dataset is shown in **boldface**, where \downarrow / \uparrow indicates that smaller/larger values of metrics are better performance.

Methods	L-UW	CCMN	PMLMD	PARd	MAE	GDF	$R_u(g)$	ComRank
scene	.398 \pm .021	.461 \pm .020	.445 \pm .032	.691 \pm .021	.424 \pm .018	.729 \pm .020	.716 \pm .026	.751\pm.015
yeast	.691 \pm .018	.650 \pm .032	.702 \pm .025	.661 \pm .039	.703 \pm .018	.714 \pm .019	.691 \pm .014	.717\pm.019
enron	.373 \pm .038	.356 \pm .052	.571 \pm .065	.432 \pm .055	.439 \pm .050	.610 \pm .078	.415 \pm .041	.626\pm.080
rcv1-s1	.450 \pm .029	.386 \pm .028	.341 \pm .034	.485 \pm .098	.422 \pm .042	.521 \pm .021	.413 \pm .044	.542\pm.099
bibtex	.234 \pm .011	.268 \pm .031	.329 \pm .104	.446 \pm .017	.285 \pm .010	.627 \pm .019	.414 \pm .018	.665\pm.021
bookmark	.525 \pm .012	.389 \pm .024	.180 \pm .005	.557 \pm .009	.518 \pm .008	.629 \pm .006	.517 \pm .013	.647\pm.006
nuswideBoW	.439 \pm .011	.402 \pm .028	.457 \pm .023	.496 \pm .014	.454 \pm .012	.556 \pm .012	.594\pm.008	.589 \pm .010

Table 4: Summary of pairwise t-test for ComRank against other comparing approaches at 0.05 significance level on biased datasets, showing in form of Win/Tie/Loss.

ComRank against	L-UW	CCMN	PMLMD	PARd	MAE	GDF	$R_u(g)$	in total
<i>One Error</i>	6/1/0	6/1/0	6/1/0	5/2/0	6/1/0	4/3/0	5/2/0	38/11/0
<i>Coverage</i>	6/1/0	7/0/0	5/2/0	7/0/0	7/0/0	4/3/0	7/0/0	43/6/0
<i>Ranking Loss</i>	6/1/0	7/0/0	5/2/0	7/0/0	7/0/0	1/6/0	7/0/0	40/9/0
<i>Average Precision</i>	7/0/0	7/0/0	5/2/0	6/1/0	6/1/0	4/3/0	6/1/0	41/8/0
in total	25/3/0	27/1/0	21/7/0	25/3/0	26/2/0	13/14/0	25/3/0	162/34/0

- PARd [Hang and Zhang, 2024]: A PML method based on a probabilistic graphical model, designed to infer potential ground-truth label information by modeling the generation process of partial multi-label data.
- MAE [Gao et al., 2023]: A MLCLL method that leverages MAE loss within the URE framework for learning.
- GDF [Gao et al., 2023]: A MLCLL method that utilizes a gradient descent-friendly loss based on URE.
- $R_u(g)$ [Gao et al., 2023]: The MLCLL loss function derived from $R_u(g)$ in Eq. (4), the URE based on uniform complementary labels.

Evaluation Metrics. We evaluate performance using four common MLL metrics: *Ranking Loss*, *Coverage*, *One Error* and *Average Precision*. Their details can be referred from Zhang and Zhou [2014]. Notably, the metric of *Ranking Loss* evaluates the fraction of reversely ordered label pairs, i.e., an irrelevant label is ranked higher than a relevant label, which is different from our method.

Implementation Details. Our experiments are conducted using PyTorch [Paszke et al., 2019] and implement on an NVIDIA TITAN RTX. To ensure fair comparisons, a linear model is applied to all datasets. For statistical analysis, we employ ten-fold cross-validation, where the dataset is randomly divided into ten subsets. The results are reported as the mean and *standard deviation* (std) of the metric. The weight decay was set to $1e-3$, and the learning rate was selected from $\{1e-3, 1e-2, 1e-1\}$. It is multiplied by 0.1 when the iteration count reaches 100 and 150 [Wang et al., 2021]. The training epochs for all datasets are 200. These settings were kept consistent across all methods.

6.2 Comparison on Uniform Complementary Labels

To evaluate the effectiveness of our method under uniform complementary label situation, we use uniform complementary-labeled data to train.

Results. Table 1 reports the results for *Average Precision*, while the results for *Coverage*, *One Error*, and *Ranking Loss* are provided in Table 8 of Appendix E due to space limitations. As shown, ComRank achieves significant improvements across all datasets. Among the 42 dataset-metric combinations, ComRank achieves the best performance in 39 cases. The most notable improvement occurs on the bibtex dataset, where *Average Precision* increases from 0.237 (under the L-UW method)




Figure 1: *Average Precision* on datasets with uniform or biased complementary labels. Dark lines show the mean of testing results, where light shadows correspond to the std.

Table 5: The running time (in 10^2 seconds) of methods with multiple uniform complementary labels.

Dataset	#Label Classes	CCMN	L-UW	GDF	MAE	$R_u(g)$	ComRank
enron	53	2.94	3.22	3.79	3.72	3.14	3.15
rcv1-s1	101	4.13	3.67	4.12	3.56	3.76	3.64
bibtex	159	5.11	3.76	4.07	3.69	3.89	3.91
bookmark	208	33.25	17.11	17.18	18.68	17.23	17.19
avg	-	11.36	6.94	7.29	7.41	7.01	6.97

to 0.658, clearly demonstrating ComRank’s strong learning capability. Additionally, compared to CLL methods, ComRank reduces the *One Error* from 0.73 to 0.308 on the enron dataset, highlighting its adaptability to MLL in the complementary label setting.

Statistical Tests. Table 2 presents the pairwise *t*-test results of ComRank against seven methods at a 0.05 significance level across four metrics. For each pairwise comparison, we use the *t*-test to determine statistical significance. If ComRank significantly outperforms the baseline, we add 1 to the win count; if ComRank is significantly worse, we add 1 to the loss count; otherwise, we add 1 to the tie count. The final results are reported in the format of win/tie/loss. ComRank consistently outperforms the baselines statistically, achieving 23/5/0 against L-UW, 25/3/0 against CCMN, and showing similar strong performance against other methods. Particularly, ComRank demonstrates dominance in metrics *One Error*, *Coverage* and *Ranking Loss*, where it achieves no losses across almost all baselines. These results highlight the effectiveness of ComRank.

Convergence Analysis on the Uniform Distribution. Figure 1 shows the epoch situation of *Average Precision* for compared methods and ComRank in bibtex and scene datasets, and similar tendency also shows on other datasets. Remaining metrics are in Figure 2 of Appendix G. As shown in the figure, ComRank demonstrates the best performance among all methods, exhibiting the fastest and most substantial improvement in *Average Precision*. Notably, GDF, $R_u(g)$, and CCMN show slight instability, with oscillations emerging midway or toward the end of training. In contrast, ComRank maintains high stability throughout the optimization process, highlighting its effectiveness in gradient descent when complementary labels are selected uniformly.

6.3 Comparison on Biased Complementary Labels

To assess the effectiveness of our method under biased distribution, we train from biased complementary-labeled data generated by the co-occurrence rates of relevant labels.

Results. Table 3 shows the results of *Average Precision*, while *Coverage*, *One Error* and *Ranking Loss* results are in Table 9 of Appendix F due to page limitation. Similar to the uniform setting, ComRank demonstrates significant improvements across all datasets. Out of the seven datasets, it achieves the best performance across all metrics on five. Among them, on the scene dataset, *One Error* is reduced dramatically from 0.852 (under L-UW) to 0.402. Compared to these methods, especially URE-based methods such as $R_u(g)$ and GDF, ComRank continues to show superior performance by reducing *Coverage* on scene from 0.422 (under $R_u(g)$) to 0.130, and improving *Average Precision* on bibtex from 0.614 (under GDF) to 0.658. These results clearly demonstrate ComRank’s advantage in handling biased complementary label scenarios.

Table 6: Comparison of Frobenius norm distances for label correlation preservation.

Datasets	CCMN	L-UW	GDF	MAE	$R_u(g)$	ComRank
scene	6.44	6.38	1.83	6.42	1.36	1.82
yeast	4.45	4.68	5.79	5.92	3.70	6.34
enron	14.92	12.88	14.33	14.84	3.33	3.65
rcv1-s1	12.42	15.06	11.32	15.04	3.20	4.95
bibtex	15.18	14.98	13.69	15.22	1.80	2.71
bookmark	13.38	15.09	11.38	15.09	1.27	11.26
nuswideBoW	14.52	14.45	9.22	14.58	4.93	8.96

Statistical Tests. Table 4 displays the pairwise t-test results of ComRank compared to various baselines on biased datasets. In particular, ComRank demonstrates superiority for all metrics with no losses in all cases. It is worth noting that, ComRank achieves near-perfect performance compared with $R_u(g)$ (25/3/0), underscoring its adaptability across datasets and biased complementary label scenarios.

Convergence Analysis on the Biased Distribution. Figure 1 shows the epoch situation of *Average Precision* for compared methods and ComRank on various datasets. Figures of *Coverage*, *Ranking Loss* and *One Error* are in Figure 3 of Appendix G. Also, the curve of ComRank shows the best performance among all methods, and remains highly stable, demonstrating its effectiveness in the gradient descent process, especially for biased complementary labels.

6.4 Further Analysis.

Complexity Analysis. With a single complementary label per instance, the proposed ComRank method has a computational complexity of $O(n(K - 1))$. Although the complexity grows with more complementary labels, our implementation avoids the high cost of pairwise comparisons by leveraging matrix operations with masking, ensuring efficiency even for large label spaces. To empirically assess scalability, we report the running time on datasets where each instance has $K/2$ uniform complementary labels (Table 5). ComRank achieves comparable or faster speeds than most baselines, demonstrating its scalability.

Label Correlation Preservation. We evaluate each method’s ability to preserve label correlations by comparing the Pearson correlation matrices of the test labels and the predicted scores, under uniform complementary labels. Their difference, measured by the Frobenius norm distance (lower is better), is reported in Table 6. ComRank achieves superior correlation preservation over most baselines, demonstrating that its ranking-based design effectively retains meaningful label dependencies.

Surrogate Loss Ablation. The table 10 in Appendix H reports the *Average Precision* from an ablation study on different surrogate losses under uniform complementary labels. We compared for log loss, sigmoid loss, softmax loss and ComRank (with exponential loss). Their details can be referred to from Appendix H. Comrank achieves competitive performance in most cases, validating its effectiveness and Bayes consistency.

7 Conclusion

In this paper, we theoretically analyze the limitations of URE, revealing that its reliance on distributional assumptions restricts its effectiveness to scenarios with uniformly selected complementary labels. Under biased complementary labels, URE struggles to provide unbiased risk estimation and fails to capture inter-label relationships. To address these issues, we propose ComRank, a complementary ranking loss framework that is theoretically justified under both uniform and biased complementary label settings. Our risk minimization analysis demonstrates that ComRank has Bayes consistency in both cases. Experimental results further validate its remarkable stability and effectiveness in learning. A current limitation of this work is that ComRank is based on multiple assumptions, including distributional assumptions and independence assumptions. In the future, we hope to extend it to more general scenarios.

References

- Peng-Fei Chen, Jun-Jie Ye, Guang-Yong Chen, Jing-Wei Zhao, and Pheng-Ann Heng. Beyond class-conditional assumption: A primary attempt to combat instance-dependent label noise. In *Proceedings of the 35th AAAI Conference on Artificial Intelligence*, volume 35, pages 11442–11450, 2021.
- Wei-Wei Cheng, Eyke Hüllermeier, and Krzysztof J Dembczynski. Bayes optimal multilabel classification via probabilistic classifier chains. In *Proceedings of the 27th ICML International Conference on Machine Learning*, pages 279–286, Haifa, Israel, 2010.
- Yu-Ting Chou, Gang Niu, Hsuan-Tien Lin, and Masashi Sugiyama. Unbiased risk estimators can mislead: A case study of learning with complementary labels. In *Proceedings of the 37th ICML International Conference on Machine Learning*, pages 1929–1938, Virtual Event, 2020.
- Lei Feng, Takuo Kaneko, Bo Han, Gang Niu, Bo An, and Masashi Sugiyama. Learning with multiple complementary labels. In *Proceedings of the 37th ICML International Conference on Machine Learning*, pages 3072–3081, Virtual Event, 2020.
- Johannes Fürnkranz, Eyke Hüllermeier, Eneldo Loza Mencía, and Klaus Brinker. Multilabel classification via calibrated label ranking. *Machine Learning*, 73:133–153, 2008.
- Wei Gao and Zhi-Hua Zhou. On the consistency of multi-label learning. *Artificial Intelligence*, 199-200:341–358, 2011.
- Yi Gao and Min-Ling Zhang. Discriminative complementary-label learning with weighted loss. In *Proceedings of the 38th ICML International Conference on Machine Learning*, pages 3587–3597, Virtual Event, 2021.
- Yi Gao, Miao Xu, and Min-Ling Zhang. Unbiased risk estimator to multi-labeled complementary label learning. In *Proceedings of the 32nd IJCAI International Joint Conference on Artificial Intelligence*, pages 3732–3740, Macao, China, 2023.
- Yi Gao, Miao Xu, and Min-Ling Zhang. Complementary to multiple labels: A correlation-aware correction approach. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 46(12): 9179–9191, 2024.
- Yi Gao, Jing-Yi Zhu, Miao Xu, and Min-Ling Zhang. Multi-label learning with multiple complementary labels. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 47(9):8013–8024, 2025.
- Jun-Yi Hang and Min-Ling Zhang. Partial multi-label learning with probabilistic graphical disambiguation. *Proceedings of the 37th NeurIPS Annual Conference on Neural Information Processing Systems*, 36:1339–1351, 2024.
- Takashi Ishida, Gang Niu, Wei-Hua Hu, and Masashi Sugiyama. Learning from complementary labels. In *Proceedings of the 30th NeurIPS Annual Conference on Neural Information Processing Systems*, pages 5639–5649, Long Beach, CA, 2017.
- Takashi Ishida, Gang Niu, Aditya Krishna Menon, and Masashi Sugiyama. Complementary-label learning for arbitrary losses and models. In *Proceedings of the 36th ICML International Conference on Machine Learning*, Proceedings of Machine Learning Research, pages 2971–2980, Long Beach, CA, 2019.
- Bin-Bin Jia, Jun-Ying Liu, Jun-Yi Hang, and Min-Ling Zhang. Learning label-specific features for decomposition-based multi-class classification. *Frontiers of Computer Science*, 17(6):176348, 2023.
- Zhi-Qiang Kou, Jing Wang, Yu-Heng Jia, Biao Liu, and Xin Geng. Instance-dependent inaccurate label distribution learning. *IEEE Transactions on Neural Networks and Learning Systems*, 36(1): 1425–1437, 2023.
- Zhi-Qiang Kou, Jing Wang, Yuheng Jia, and Xin Geng. Inaccurate label distribution learning. *IEEE Transactions on Circuits and Systems for Video Technology*, 34(10):10237–10249, 2024.

- Jia-Xuan Li, Xiao-Yan Zhu, Wei-Chu Zhang, and Jia-Yin Wang. A ranking-based problem transformation method for weakly supervised multi-label learning. *Pattern Recognition*, 153:110505, 2024.
- Yun-Cheng Li, Yale Song, and Jie-Bo Luo. Improving pairwise ranking for multi-label image classification. In *Proceedings of the 35th CVPR Conference on Computer Vision and Pattern Recognition*, pages 1837–1845, Honolulu, HI, 2017.
- Yi Liu, Rong Jin, and Liu Yang. Semi-supervised multi-label learning by constrained non-negative matrix factorization. In *Proceedings of the 20th AAAI Conference on Artificial Intelligence*, volume 6, pages 421–426, Boston, USA, 2006.
- Xue-Song Niu, Hu Han, Shi-Guang Shan, and Xi-Lin Chen. Multi-label co-regularization for semi-supervised facial action unit recognition. In *Proceedings of the 32th NeurIPS Annual Conference on Neural Information Processing Systems*, pages 907–917, Vancouver, Canada, 2019.
- Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Ze-Ming Lin, Natalia Gimelshein, Luca Antiga, Alban Desmaison, Andreas Köpf, Edward Yang, Zachary DeVito, Martin Raison, Alykhan Tejani, Sasank Chilamkurthy, Benoit Steiner, Lu Fang, Jun-Jie Bai, and Soumith Chintala. Pytorch: An imperative style, high-performance deep learning library. In *Proceedings of the 32th NeurIPS Annual Conference on Neural Information Processing Systems*, pages 8024–8035, Vancouver, Canada, 2019.
- Jiang-Xin Shi, Tong Wei, and Yu-Feng Li. Residual diverse ensemble for long-tailed multi-label text classification. *Science China Information Sciences*, 67(11):212102, 2024.
- Yu-Yin Sun, Yin Zhang, and Zhi-Hua Zhou. Multi-label learning with weak label. In *Proceedings of the 24th AAAI Conference on Artificial Intelligence*, volume 24, pages 593–598, Atlanta, USA, 2010.
- Yi Tang, Yi Gao, Yong-Gang Luo, Ju-Cheng Yang, Miao Xu, and Min-Ling Zhang. Unlearning from weakly supervised learning. In *Proceedings of the 32nd IJCAI International Joint Conference on Artificial Intelligence*, pages 5000–5008, Macao, China, 2023.
- Deng-Bao Wang, Lei Feng, and Min-Ling Zhang. Learning from complementary labels via partial-output consistency regularization. In *Proceedings of the 30th IJCAI International Joint Conference on Artificial Intelligence*, pages 3075–3081, Virtual Event, 2021.
- Bao-Yuan Wu, Zhi-Lei Liu, Shang-Fei Wang, Bao-Gang Hu, and Qiang Ji. Multi-label learning with missing labels. In *Proceedings of the 22nd ICPR International Conference on Pattern Recognition*, pages 1964–1968, Stockholm, Sweden, 2014.
- Xiaobo Xia, Tongliang Liu, Bo Han, Nannan Wang, Mingming Gong, Haifeng Liu, Gang Niu, Dacheng Tao, and Masashi Sugiyama. Part-dependent label noise: Towards instance-dependent label noise. *Proceedings of the 33th NeurIPS Annual Conference on Neural Information Processing Systems*, 33:7597–7610, 2020.
- Ming-Kun Xie and Sheng-Jun Huang. Partial multi-label learning. In *Proceedings of the 32th AAAI Conference on Artificial Intelligence*, volume 32, page 4302–4309, 2018.
- Ming-Kun Xie and Sheng-Jun Huang. CCMN: A general framework for learning with class-conditional multi-label noise. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 45(1):154–166, 2023.
- Ming-Kun Xie, Feng Sun, and Sheng-Jun Huang. Partial multi-label learning with meta disambiguation. In *Proceedings of the 27th ACM SIGKDD Conference on Knowledge Discovery & Data Mining*, pages 1904–1912, 2021.
- Xi-Yu Yu, Tong-Liang Liu, Ming-Ming Gong, and Da-Cheng Tao. Learning with biased complementary labels. In *Proceedings of the 15th ECCV European Conference on Computer Vision*, pages 68–83, 2018.
- Min-Ling Zhang and Jun-Peng Fang. Partial multi-label learning via credible label elicitation. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 43(10):3587–3599, 2020.

Min-Ling Zhang and Zhi-Hua Zhou. A review on multi-label learning algorithms. *IEEE Transactions on Knowledge and Data Engineering*, 26(8):1819–1837, 2014.

Min-Ling Zhang, Yu-Kun Li, Xu-Ying Liu, and Xin Geng. Binary relevance for multi-label learning: an overview. *Frontiers Computer Science*, 12(2):191–202, 2018.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

Answer: [\[Yes\]](#)

Justification: Please see the abstract and introduction sections.

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: Please see section 7.

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [\[Yes\]](#)

Justification: For Theorem 4.5, it has Assumption 5, Assumption 4.4 and the proof is stated in Appendix B; For Lemma 5.1, the proof is stated in Appendix C; For Theorem 5.2, the proof is stated in Appendix D;

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [\[Yes\]](#)

Justification: The detailed method and experimental settings are provided in the paper, which enable the proposed method to be reproduced with public datasets.

Guidelines:

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [No]

Justification: The datasets used in this paper are public, and their download links are offered in the paper. The code for this paper will be released after the paper is accepted.

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: Experimental Setting can be found in the section 6.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Please see section 6.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).

- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Please see section 6.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The research conducted in the paper conform to the NeurIPS Code of Ethics.

Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Please see section 1.

Guidelines:

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.

- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [NA]

Justification: This paper does not include the above content.

Guidelines:

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: The datasets used in this paper are public.

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.

- If this information is not available online, the authors are encouraged to reach out to the asset’s creators.

13. **New assets**

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [NA]

Justification: The paper does not release new assets.

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. **Crowdsourcing and research with human subjects**

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [NA]

Justification: The paper does not involve crowdsourcing nor research with human subjects.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [NA]

Justification: The paper does not involve crowdsourcing nor research with human subjects.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. **Declaration of LLM usage**

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does not impact the core methodology, scientific rigorousness, or originality of the research, declaration is not required.

Answer: [NA]

Justification: The core method development in this research does not involve LLMs as any important, original, or non-standard components.

Guidelines:

- The answer NA means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy (<https://neurips.cc/Conferences/2025/LLM>) for what should or should not be described.

Appendix

A Details of Datasets

Table 7 describes the datasets used in the experiments of this paper, where #Instance represents the number of samples, #Features represents the sample feature dimension, #Label Classes represents the label space, #Cardinality represents the average number of relevant labels of the sample, and #Domain represents the data type.

Table 7: Characteristics of datasets.

Datasets	#Instances	#Features	#Label Classes	#Cardinality	#Domain
scene	2407	294	6	1.07	images
yeast	2417	103	14	4.23	biology
enron	1702	1001	53	3.38	Text
rcv1-s1	5815	944	101	2.88	Text
bibtex	7365	1836	159	2.40	images
bookmark	87856	2150	208	1.25	Text
nuswideBoW	269648	500	81	1.87	images

B The Proof of Theorem 4.5

Theorem 4.5. *URE under biased distribution: Given $p(k = \bar{y}|\mathbf{x})$ and $R(\mathbf{g})$, the equality $R(\mathbf{g}) = \bar{R}(\mathbf{g}) = R_b(\mathbf{g})$ holds when*

$$R_b(\mathbf{g}) = \mathbb{E}_{p(\mathbf{x}, \bar{y})} \left[\sum_{Y \subseteq \mathcal{Y}, \bar{y} \neq Y} \bar{l}_{Y\bar{y}}^+ \mathcal{L}(\mathbf{g}, Y) \right], \quad (7)$$

where $\bar{l}_{Y\bar{y}}^+$ belongs to the matrix $\bar{\mathbf{L}}^+ = [\bar{l}^+]_{|\mathcal{Y}| \times K}$, and $\bar{\mathbf{L}}^+$ is the Moore-Penrose pseudoinverse of $\bar{\mathbf{L}}^T$.

Proof. According to Assumption 4.3 and Assumption 4.4, for $\bar{y} \notin Y$, $p(\bar{y}|\mathbf{x}, Y) = \frac{z}{l_{Y\bar{y}}}$. Therefore,

$$\begin{aligned} p(\mathbf{x}, \bar{y}) &= \sum_{Y \subseteq \mathcal{Y}, \bar{y} \notin Y} p(\mathbf{x}, Y, \bar{y}) = \sum_{Y \subseteq \mathcal{Y}, \bar{y} \notin Y} p(\mathbf{x}, Y) p(\bar{y}|\mathbf{x}, Y) \\ &= \sum_{Y \subseteq \mathcal{Y}, \bar{y} \notin Y} \frac{z}{l_{Y\bar{y}}} p(\mathbf{x}, Y). \end{aligned}$$

Set $\bar{\mathbf{L}} = [\bar{l}]_{|\mathcal{Y}| \times K}$, where $\bar{l}_{Yk} = \begin{cases} \frac{z}{l_{Yk}}, & k \in Y \\ 0, & k \notin Y \end{cases}$. By expanding $p(\mathbf{x}, \bar{y})$ and $p(\mathbf{x}, Y)$ with respect to \bar{y} and Y as marginal probabilities, we can obtain:

$$\begin{bmatrix} p(\mathbf{x}, \bar{y} = 1) \\ \vdots \\ p(\mathbf{x}, \bar{y} = k) \\ \vdots \\ p(\mathbf{x}, \bar{y} = K) \end{bmatrix}_{K \times 1} = \bar{\mathbf{L}}^T \begin{bmatrix} p(\mathbf{x}, Y = Y_1) \\ p(\mathbf{x}, Y = Y_2) \\ \vdots \\ p(\mathbf{x}, Y = Y_{|\mathcal{Y}|}) \end{bmatrix}_{|\mathcal{Y}| \times 1}.$$

When $\bar{\mathbf{L}}^T$ is full rank, there exists a matrix $\bar{\mathbf{L}}^+ = [\bar{l}^+]_{|\mathcal{Y}| \times K}$, which is the Moore-Penrose pseudoinverse of $\bar{\mathbf{L}}^T$, satisfies:

$$\bar{\mathbf{L}}^+ \cdot \begin{bmatrix} p(\mathbf{x}, \bar{y} = 1) \\ \vdots \\ p(\mathbf{x}, \bar{y} = k) \\ \vdots \\ p(\mathbf{x}, \bar{y} = K) \end{bmatrix}_{K \times 1} = \begin{bmatrix} p(\mathbf{x}, Y = Y_1) \\ p(\mathbf{x}, Y = Y_2) \\ \vdots \\ p(\mathbf{x}, Y = Y_{|\mathcal{Y}|}) \end{bmatrix}_{|\mathcal{Y}| \times 1}.$$

Therefore, we can have the relationship

$$p(\mathbf{x}, Y) = \sum_{\bar{y}=1, \bar{y} \neq Y}^K \bar{l}_{Y\bar{y}}^+ p(\mathbf{x}, \bar{y}),$$

where $\bar{l}_{Y\bar{y}}^+$ is the \bar{y} -th column of the row corresponding to Y in $\bar{\mathbf{L}}^+$. Accordingly, the URE under Assumption 4.3 is:

$$\begin{aligned} R_b(\mathbf{g}) &= \mathbb{E}_{p(\mathbf{x}, Y)} [\mathcal{L}(\mathbf{g}, Y)] \\ &= \int_{\mathcal{X}} \sum_{Y \subseteq \mathcal{Y}} \mathcal{L}(\mathbf{g}, Y) p(\mathbf{x}, Y) d\mathbf{x} \\ &= \int_{\mathcal{X}} \sum_{Y \subseteq \mathcal{Y}} \sum_{\bar{y}=1}^K \mathcal{L}(\mathbf{g}, Y) \bar{l}_{Y\bar{y}}^+ p(\mathbf{x}, \bar{y}) d\mathbf{x} \\ &= \int_{\mathcal{X}} \sum_{\bar{y}=1}^K \sum_{Y \subseteq \mathcal{Y}, \bar{y} \neq Y} \mathcal{L}(\mathbf{g}, Y) \bar{l}_{Y\bar{y}}^+ \bar{p}(\mathbf{x}, \bar{y}) d\mathbf{x} \\ &= \mathbb{E}_{p(\mathbf{x}, \bar{y})} \left[\sum_{Y \subseteq \mathcal{Y}, \bar{y} \neq Y} \bar{l}_{Y\bar{y}}^+ \mathcal{L}(\mathbf{g}, Y) \right]. \end{aligned}$$

□

C The Proof of Lemma 5.1

Lemma 5.1. Under $\bar{\mathcal{L}}_{\text{CR}}(\mathbf{g}(\mathbf{x}), \bar{y})$, $\bar{\mathbf{g}}^k(\mathbf{x}) \geq \bar{\mathbf{g}}^j(\mathbf{x})$ if and only if $p(k = \bar{y}|\mathbf{x}) \leq p(j = \bar{y}|\mathbf{x})$.

Proof. Let $\Delta Y_k = \{0, 1, -1\} \in \mathbb{R}^{1 \times K}$, where the \bar{y} -th position is 1, the k -th position is -1, and other positions are 0, then:

$$\bar{\mathcal{L}}_{\text{CR}}(\mathbf{g}(\mathbf{x}), \bar{y}) = \sum_{k=1, k \neq \bar{y}}^K \exp(\mathbf{g}(\mathbf{x}) \Delta Y_k^T).$$

Assume $\alpha_k = \exp(\mathbf{g}(\mathbf{x}) \Delta Y_k^T)$, we can obtain

$$\begin{aligned} \bar{R}(\mathbf{g}) &= \mathbb{E}_p(\mathbf{x}, \bar{y}) [\bar{\mathcal{L}}_{\text{CR}}(\mathbf{g}(\mathbf{x}), \bar{y})] \\ &= \int_{\mathcal{X}} \int_{\mathcal{Y}} \bar{\mathcal{L}}_{\text{CR}}(\mathbf{g}, \bar{y}) p(\mathbf{x}, \bar{y}) d\mathbf{x} d\bar{y} \\ &= \int_{\mathcal{X}} \sum_{\bar{y} \in \mathcal{Y}} \bar{\mathcal{L}}_{\text{CR}}(\mathbf{g}, \bar{y}) p(\mathbf{x}, \bar{y}) d\mathbf{x} \\ &= \int_{\mathcal{X}} \sum_{\bar{y} \in \mathcal{Y}} \sum_{k=1, k \neq \bar{y}}^K \sum_{j=1, j = \bar{y}}^K \alpha_k p(\mathbf{x}, \bar{y}) d\mathbf{x} \\ &= \int_{\mathcal{X}} \sum_{k=1}^K \sum_{j=1}^K \sum_{\bar{y} \in \mathcal{Y}, k \neq \bar{y}, j = \bar{y}} \alpha_k p(\mathbf{x}, \bar{y}) d\mathbf{x} \end{aligned}$$

$$\begin{aligned}
&= \int_{\mathcal{X}} \sum_{k=1}^K \sum_{j=1}^K \sum_{\bar{y} \in \mathcal{Y}, k \neq \bar{y}, j = \bar{y}} \alpha_k p(\bar{y}|\mathbf{x}) p(\mathbf{x}) d\mathbf{x} \\
&= \int_{\mathcal{X}} \sum_{k=1}^K \sum_{j=1}^K p(k \neq \bar{y}, j = \bar{y}|\mathbf{x}) \alpha_k p(\mathbf{x}) d\mathbf{x}.
\end{aligned}$$

Since the minimization is performed with respect to \mathbf{g} , and \mathbf{x} does not affect the minimization, $\bar{\mathbf{g}}^* = \arg \min_{\mathbf{g} \in G} \bar{R}(\mathbf{g})$ can be converted to

$$\bar{\mathbf{g}}'^* = \arg \min_{\mathbf{g} \in G} \bar{R}'(\mathbf{g}) = \sum_{k=1}^K \sum_{j=1}^K p(k \neq \bar{y}, j = \bar{y}|\mathbf{x}) \alpha_k$$

When $\bar{R}'(\mathbf{g})$ reaches $\mathbf{g}'^*(\mathbf{x})$, $\bar{R}(\mathbf{g})$ can reach $\mathbf{g}^*(\mathbf{x})$. Let $\beta_k = p(k \neq \bar{y}, j = \bar{y}|\mathbf{x})$, thus the first-order derivative is

$$\frac{\partial \bar{R}'}{\partial \mathbf{g}} = \sum_{k=1}^K \sum_{j=1}^K \beta_k \Delta Y_k \alpha_k.$$

And the second derivative is

$$\frac{\partial^2 \bar{R}'}{\partial \mathbf{g}^2} = \sum_{k=1}^K \sum_{j=1}^K \beta_k \Delta Y_k^T \Delta Y_k \alpha_k \geq \mathbf{0}_{K \times K}.$$

Therefore when the first derivative is equal to 0, $\bar{R}'(\mathbf{g})$ has the minimum. Let $\frac{\partial \bar{R}'}{\partial \mathbf{g}} = \mathbf{0}$, we have

$$\begin{aligned}
&\sum_{k=1}^K \sum_{j=1}^K \beta_k \Delta Y_k \alpha_k = \mathbf{0} \\
&\Rightarrow \sum_{k=1}^K \sum_{j=1}^K \beta_k \Delta Y_k \exp(\bar{\mathbf{g}}'^*(\mathbf{x}) \Delta Y_k^T) = \mathbf{0} \\
&\Rightarrow \sum_{k=1}^K \sum_{j=1}^K \beta_k \Delta Y_k \exp(\bar{\mathbf{g}}'_j(\mathbf{x}) - \bar{\mathbf{g}}'_k(\mathbf{x})) = \mathbf{0}. \tag{8}
\end{aligned}$$

Therefore, the first-order derivative is $\mathbf{0}$ when

$$\exp(\bar{\mathbf{g}}'_j(\mathbf{x}) - \bar{\mathbf{g}}'_k(\mathbf{x})) = \frac{p(k = \bar{y}, j \neq \bar{y}|\mathbf{x})}{p(k \neq \bar{y}, j = \bar{y}|\mathbf{x})},$$

because at this moment for each dimension in Eq. (8),

$$\sum_{j=1}^K p(k = \bar{y}, j \neq \bar{y}|\mathbf{x}) - \sum_{k=1}^K p(k \neq \bar{y}, j = \bar{y}|\mathbf{x}) = 0.$$

Then, we have

$$\begin{aligned}
\bar{\mathbf{g}}'_j(\mathbf{x}) - \bar{\mathbf{g}}'_k(\mathbf{x}) &= \log \frac{p(k = \bar{y}, j \neq \bar{y}|\mathbf{x})}{p(k \neq \bar{y}, j = \bar{y}|\mathbf{x})}, \quad \forall k, j \in \mathcal{Y} \\
\Rightarrow \bar{\mathbf{g}}'_k(\mathbf{x}) - \bar{\mathbf{g}}'_j(\mathbf{x}) &= \log \frac{p(k \neq \bar{y}, j = \bar{y}|\mathbf{x})}{p(k = \bar{y}, j \neq \bar{y}|\mathbf{x})}, \quad \forall k, j \in \mathcal{Y}.
\end{aligned}$$

Therefore, $\bar{\mathbf{g}}'_k(\mathbf{x}) > \bar{\mathbf{g}}'_j(\mathbf{x})$ if and only if $p(k \neq \bar{y}, j = \bar{y}|\mathbf{x}) \geq p(k = \bar{y}, j \neq \bar{y}|\mathbf{x})$ holds. That is, $\bar{\mathbf{g}}_k^*(\mathbf{x}) > \bar{\mathbf{g}}_j^*(\mathbf{x})$ if and only if $p(k \neq \bar{y}, j = \bar{y}|\mathbf{x}) \geq p(k = \bar{y}, j \neq \bar{y}|\mathbf{x})$ holds. Then, we have

$$\begin{aligned}
&p(k \neq \bar{y}|j = \bar{y}, \mathbf{x}) p(j = \bar{y}|\mathbf{x}) \geq p(j \neq \bar{y}|k = \bar{y}, \mathbf{x}) p(k = \bar{y}|\mathbf{x}) \\
&\quad \because p(k \neq \bar{y}|j = \bar{y}, \mathbf{x}) = p(k \neq \bar{y}|\mathbf{x}) - p(k \neq \bar{y}, j \neq \bar{y}|\mathbf{x}) \\
&\Rightarrow [p(k \neq \bar{y}|\mathbf{x}) - p(k \neq \bar{y}, j \neq \bar{y}|\mathbf{x})] p(j = \bar{y}|\mathbf{x}) \geq [p(j \neq \bar{y}|\mathbf{x}) - p(j \neq \bar{y}, k \neq \bar{y}|\mathbf{x})] p(k = \bar{y}|\mathbf{x})
\end{aligned}$$

$$\begin{aligned}
&\Rightarrow p(k \neq \bar{y}|\mathbf{x})p(j = \bar{y}|\mathbf{x}) - p(k \neq \bar{y}, j \neq \bar{y}|\mathbf{x})p(j = \bar{y}|\mathbf{x}) \\
&\geq p(j \neq \bar{y}|\mathbf{x})p(k = \bar{y}|\mathbf{x}) - p(j \neq \bar{y}, k \neq \bar{y}|\mathbf{x})p(k = \bar{y}|\mathbf{x}) \\
&\because p(j = \bar{y}|\mathbf{x}) = 1 - p(j \neq \bar{y}|\mathbf{x}), p(k = \bar{y}|\mathbf{x}) = 1 - p(k \neq \bar{y}|\mathbf{x}) \\
&\Rightarrow p(k \neq \bar{y}|\mathbf{x}) - p(k \neq \bar{y}|\mathbf{x})p(j \neq \bar{y}|\mathbf{x}) - p(k \neq \bar{y}, j \neq \bar{y}|\mathbf{x}) + p(k \neq \bar{y}, j \neq \bar{y}|\mathbf{x})p(j \neq \bar{y}|\mathbf{x}) \\
&\geq p(j \neq \bar{y}|\mathbf{x}) - p(j \neq \bar{y}|\mathbf{x})p(k \neq \bar{y}|\mathbf{x}) - p(j \neq \bar{y}, k \neq \bar{y}|\mathbf{x}) + p(j \neq \bar{y}, k \neq \bar{y}|\mathbf{x})p(k \neq \bar{y}|\mathbf{x}) \\
&\Rightarrow p(k \neq \bar{y}|\mathbf{x}) - p(j \neq \bar{y}|\mathbf{x}) \geq p(k \neq \bar{y}, j \neq \bar{y}|\mathbf{x})(p(k \neq \bar{y}|\mathbf{x}) - p(j \neq \bar{y}|\mathbf{x})).
\end{aligned}$$

This means if and only if $p(k \neq \bar{y}) \geq p(j \neq \bar{y}|\mathbf{x})$ does the inequality hold. Therefore, $\bar{g}_k^*(\mathbf{x}) \geq \bar{g}_j^*(\mathbf{x})$ if and only if $p(k \neq \bar{y}|\mathbf{x}) \geq p(j \neq \bar{y}|\mathbf{x})$ holds. \square

D The Proof of Theorem 5.2

Theorem 5.2. *Bayes Consistency for ComRank: For both uniform and biased complementary label distributions, $\bar{g}^k(\mathbf{x}) \geq \bar{g}^j(\mathbf{x})$ holds under $\bar{\mathcal{L}}_{\text{CR}}(\mathbf{g}(\mathbf{x}), \bar{y})$ if and only if $p(k \in Y|\mathbf{x}) \geq p(j \in Y|\mathbf{x})$.*

Theorem 5.2.1. *Under $\bar{\mathcal{L}}_{\text{CR}}(\mathbf{g}(\mathbf{x}), \bar{y})$, $\bar{g}^k(\mathbf{x}) \geq \bar{g}^j(\mathbf{x})$ if and only if $p(k \in Y|\mathbf{x}) \geq p(j \in Y|\mathbf{x})$, for Assumption 4.1.*

Proof. According to Lemma 5.1, $\bar{g}_k^*(\mathbf{x}) > \bar{g}_j^*(\mathbf{x})$ if and only if $p(k = \bar{y}|\mathbf{x}) \leq p(j = \bar{y}|\mathbf{x})$.

Based on Assumption 4.1, whenever $k \in \bar{y}$ and $j \notin \bar{y}$, it holds that $p(k = \bar{y}|\mathbf{x}) \leq p(j = \bar{y}|\mathbf{x})$.

Additionally, when $k \in \bar{y}$ and $j \in \bar{y}$, $p(k \in \bar{y}|\mathbf{x}) \geq p(j \in \bar{y}|\mathbf{x})$.

That is, $\bar{g}_k^*(\mathbf{x}) > \bar{g}_j^*(\mathbf{x})$ if and only if $p(k \in \bar{y}|\mathbf{x}) \geq p(j \in \bar{y}|\mathbf{x})$, which satisfies Bayes consistency. \square

Theorem 5.2.2. *Under $\bar{\mathcal{L}}_{\text{CR}}(\mathbf{g}(\mathbf{x}), \bar{y})$, $\bar{g}^k(\mathbf{x}) \geq \bar{g}^j(\mathbf{x})$ if and only if $p(k \in Y|\mathbf{x}) \geq p(j \in Y|\mathbf{x})$, for Assumption 4.3.*

Proof. According to Lemma 5.1, $\bar{g}_k^*(\mathbf{x}) > \bar{g}_j^*(\mathbf{x})$ if and only if $p(k = \bar{y}|\mathbf{x}) \leq p(j = \bar{y}|\mathbf{x})$. According to Assumption 4.3, when $p(k = \bar{y}|\mathbf{x}) \leq p(j = \bar{y}|\mathbf{x})$, there are two possible cases:

1. $k \in Y, j \notin Y$: In this case, we must have $p(k \in Y|\mathbf{x}) \geq p(j \in Y|\mathbf{x})$.

That is, $\bar{g}_k^*(\mathbf{x}) \geq \bar{g}_j^*(\mathbf{x})$ if and only if $p(k \in Y|\mathbf{x}) \geq p(j \in Y|\mathbf{x})$, which satisfies Bayes consistency.

2. $k, j \notin Y$: Since $p(k = \bar{y}|\mathbf{x}) = \frac{z}{l_{Yk}}$, we obtain $\Rightarrow \frac{z}{l_{Yk}} \leq \frac{z}{l_{Yj}}$.

Since z is same, we have $l_{Yk} \geq l_{Yj}$.

$\because l_{Yk} \propto p(k \in Y|\mathbf{x})$, we have $p(k \in Y|\mathbf{x}) \geq p(j \in Y|\mathbf{x})$.

That is, $\bar{g}_k^*(\mathbf{x}) > \bar{g}_j^*(\mathbf{x})$ if and only if $p(k \in Y|\mathbf{x}) \geq p(j \in Y|\mathbf{x})$, which satisfies Bayes consistency. \square

E Comparison on Uniform Complementary Labels

Table 8 shows the comparison of ComRank against multiple methods in *One Error*, *Ranking Loss* and *Coverage* under uniform complementary labels. As we can see, ComRank demonstrates strong performance across all methods and achieves superior results on most datasets.

F Comparison on Biased Complementary Labels

Table 9 shows the comparison of ComRank against multiple methods in *One Error*, *Ranking Loss* and *Coverage* under biased complementary labels. As we can see, ComRank shows competitive results among all methods and outperforms most datasets.

Table 8: Experimental results (mean \pm std) on the training data with uniform complementary labels. The best performance of each dataset is shown in **boldface**, where \downarrow / \uparrow indicates that smaller/larger values of metric are better performance.

Methods	L-UW	CCMN	PMLMD	PARD	MAE	GDF	$R_u(g)$	ComRank
<i>Coverage\downarrow</i>								
scene	.437 \pm .022	.404 \pm .021	.388 \pm .033	.189 \pm .020	.412 \pm .020	.157 \pm .016	.172 \pm .021	.143\pm.013
yeast	.565 \pm .016	.664 \pm .068	.509 \pm .033	.622 \pm .066	.548 \pm .017	.534 \pm .015	.569 \pm .021	.520\pm.019
enron	.633 \pm .059	.652 \pm .030	.477 \pm .106	.602 \pm .052	.589 \pm .060	.470 \pm .049	.602 \pm .060	.458\pm.051
rcv1-s1	.343 \pm .027	.472 \pm .042	.537 \pm .034	.414 \pm .087	.381 \pm .036	.297 \pm .028	.395 \pm .014	.305\pm.060
bibtex	.499 \pm .014	.463 \pm .033	.424 \pm .104	.299 \pm .013	.438 \pm .016	.218 \pm .017	.322 \pm .018	.195\pm.015
bookmark	.289 \pm .007	.395 \pm .022	.669 \pm .028	.275 \pm .007	.299 \pm .007	.210 \pm .005	.294 \pm .009	.199\pm.005
nuswideBoW	.425 \pm .010	.473 \pm .037	.385 \pm .036	.369 \pm .011	.409 \pm .010	.313 \pm .011	.309 \pm .012	.298\pm.011
<i>Ranking Loss\downarrow</i>								
scene	.523 \pm .021	.466 \pm .023	.467 \pm .048	.172 \pm .033	.488 \pm .023	.153 \pm .009	.175 \pm .011	.136\pm.011
yeast	.245 \pm .013	.294 \pm .022	.221 \pm .018	.355 \pm .115	.231 \pm .014	.219 \pm .014	.250 \pm .015	.215\pm.013
enron	.438 \pm .028	.481 \pm .034	.262 \pm .030	.370 \pm .031	.387 \pm .034	.231 \pm .029	.395 \pm .024	.221\pm.029
rcv1-s1	.272 \pm .032	.344 \pm .044	.447 \pm .044	.299 \pm .048	.307 \pm .030	.275 \pm .058	.369 \pm .017	.267\pm.081
bibtex	.491 \pm .009	.479 \pm .031	.460 \pm .102	.280 \pm .013	.426 \pm .009	.206 \pm .013	.308 \pm .014	.180\pm.012
bookmark	.280 \pm .007	.368 \pm .032	.661 \pm .024	.269 \pm .009	.287 \pm .006	.196 \pm .005	.280 \pm .008	.187\pm.007
nuswideBoW	.287 \pm .005	.340 \pm .018	.274 \pm .026	.317 \pm .094	.274 \pm .005	.206 \pm .005	.196 \pm .006	.192\pm.006
<i>One Error\downarrow</i>								
scene	.851 \pm .017	.758 \pm .035	.792 \pm .057	.411 \pm .021	.803 \pm .024	.384 \pm .020	.418 \pm .022	.352\pm.014
yeast	.252 \pm .024	.261 \pm .035	.270 \pm .038	.324 \pm .202	.251 \pm .023	.249 \pm .025	.277 \pm .017	.249\pm.025
enron	.723 \pm .060	.781 \pm .058	.519 \pm .136	.588 \pm .074	.647 \pm .075	.330 \pm .116	.668 \pm .074	.308\pm.124
rcv1-s1	.708 \pm .024	.708 \pm .059	.750 \pm .072	.624 \pm .144	.702 \pm .043	.652 \pm .085	.777 \pm .031	.609\pm.174
bibtex	.915 \pm .010	.894 \pm .024	.881 \pm .033	.720 \pm .025	.871 \pm .015	.480 \pm .023	.744 \pm .018	.418\pm.021
bookmark	.617 \pm .011	.745 \pm .040	.927 \pm .016	.550 \pm .011	.600 \pm .011	.483 \pm .008	.608 \pm .007	.474\pm.006
nuswideBoW	.643 \pm .019	.710 \pm .040	.690 \pm .066	.630 \pm .044	.629 \pm .023	.540 \pm .023	.488\pm.017	.500 \pm .024

Table 9: Experimental results (mean \pm std) on the training data with biased complementary labels. The best performance of each dataset is shown in **boldface**, where \downarrow / \uparrow indicates that smaller/larger values of metric are better performance.

Methods	L-UW	CCMN	PMLMD	PARD	MAE	GDF	$R_u(g)$	ComRank
<i>Coverage\downarrow</i>								
scene	.450 \pm .018	.402 \pm .019	.405 \pm .039	.160 \pm .028	.422 \pm .020	.144 \pm .008	.162 \pm .009	.130\pm.010
yeast	.575 \pm .018	.648 \pm .055	.514\pm.029	.712 \pm .081	.557 \pm .019	.540 \pm .021	.582 \pm .025	.530 \pm .019
enron	.634 \pm .059	.681 \pm .036	.494 \pm .040	.587 \pm .067	.600 \pm .062	.467 \pm .060	.603 \pm .060	.456\pm.058
rcv1-s1	.348 \pm .049	.425 \pm .038	.524 \pm .049	.397 \pm .047	.384 \pm .042	.354 \pm .050	.441 \pm .008	.350\pm.061
bibtex	.497 \pm .012	.485 \pm .026	.463 \pm .103	.300 \pm .011	.437 \pm .012	.230 \pm .015	.327 \pm .017	.206\pm.014
bookmark	.299 \pm .007	.381 \pm .029	.655 \pm .024	.289 \pm .009	.305 \pm .007	.219 \pm .006	.300 \pm .008	.210\pm.008
nuswideBoW	.404 \pm .011	.449 \pm .020	.384 \pm .037	.451 \pm .115	.391 \pm .011	.313 \pm .011	.307 \pm .012	.298\pm.011
<i>Ranking Loss\downarrow</i>								
scene	.506 \pm .025	.466 \pm .025	.447 \pm .039	.208 \pm .022	.476 \pm .024	.171 \pm .016	.188 \pm .024	.154\pm.014
yeast	.240 \pm .011	.298 \pm .040	.216 \pm .026	.277 \pm .034	.228 \pm .011	.218 \pm .011	.241 \pm .011	.212\pm.011
enron	.439 \pm .029	.452 \pm .043	.250 \pm .055	.386 \pm .039	.377 \pm .034	.233 \pm .029	.396 \pm .030	.221\pm.024
rcv1-s1	.271 \pm .018	.383 \pm .033	.452 \pm .036	.321 \pm .103	.305 \pm .027	.223 \pm .035	.323 \pm .026	.230\pm.070
bibtex	.493 \pm .012	.455 \pm .034	.423 \pm .115	.278 \pm .011	.427 \pm .013	.194 \pm .015	.303 \pm .016	.170\pm.013
bookmark	.270 \pm .007	.386 \pm .025	.677 \pm .028	.255 \pm .007	.281 \pm .007	.186 \pm .005	.274 \pm .009	.175\pm.005
nuswideBoW	.309 \pm .005	.363 \pm .040	.274 \pm .024	.256 \pm .010	.293 \pm .004	.208 \pm .006	.200 \pm .006	.193\pm.006
<i>One Error\downarrow</i>								
scene	.852 \pm .026	.749 \pm .033	.807 \pm .052	.482 \pm .031	.820 \pm .018	.435 \pm .030	.447 \pm .036	.402\pm.024
yeast	.254 \pm .024	.254 \pm .029	.267 \pm .034	.293 \pm .127	.252 \pm .025	.253\pm.024	.265 \pm .025	.255 \pm .025
enron	.727 \pm .058	.764 \pm .073	.493 \pm .108	.632 \pm .104	.627 \pm .079	.341 \pm .136	.669 \pm .068	.324\pm.144
rcv1-s1	.693 \pm .032	.704 \pm .080	.737 \pm .085	.568 \pm .125	.716 \pm .053	.611 \pm .038	.724 \pm .068	.561\pm.150
bibtex	.917 \pm .011	.886 \pm .033	.811 \pm .126	.717 \pm .024	.873 \pm .014	.461 \pm .025	.747 \pm .019	.414\pm.029
bookmark	.589 \pm .016	.755 \pm .040	.926 \pm .013	.546 \pm .011	.582 \pm .011	.469 \pm .008	.604 \pm .016	.448\pm.009
nuswideBoW	.649 \pm .018	.719 \pm .058	.668 \pm .055	.608 \pm .026	.636 \pm .021	.537 \pm .027	.490\pm.017	.496 \pm .023

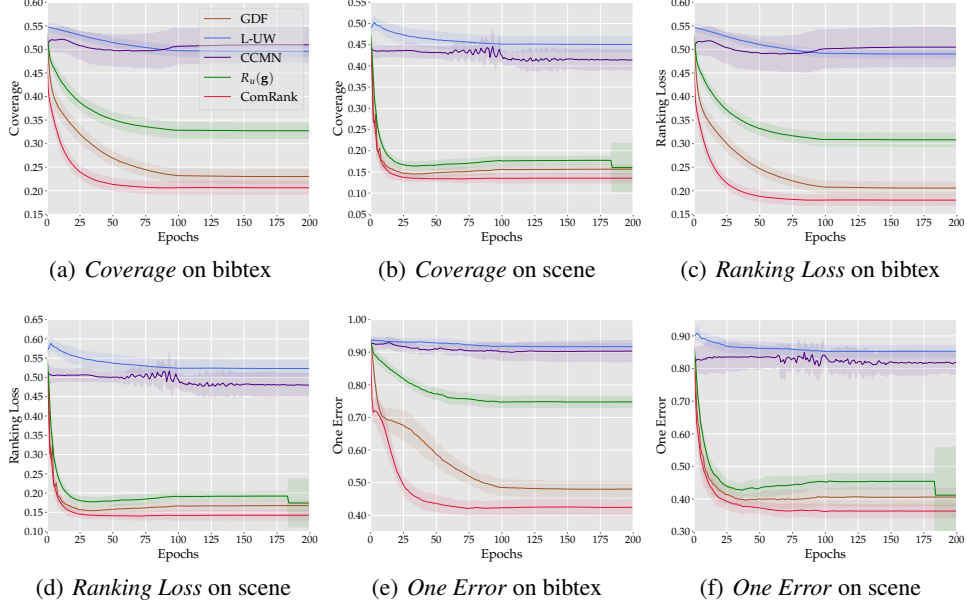


Figure 2: Coverage, Ranking Loss and One Error on various datasets with uniform complementary labels.

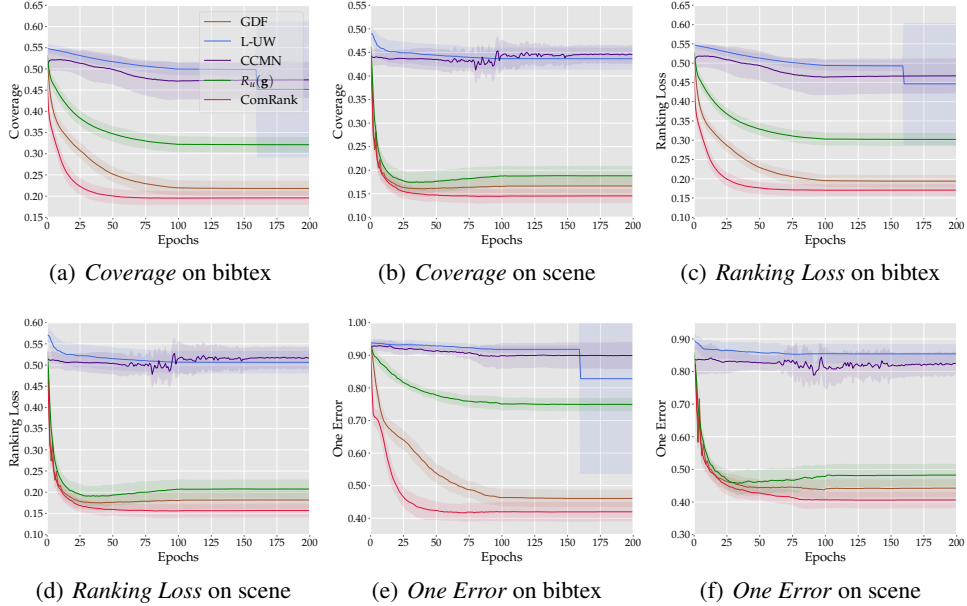


Figure 3: Coverage, Ranking Loss and One Error on various datasets with biased complementary labels. Dark lines show the mean of testing results, where light shadows correspond to the std.

G Figure

Figure 2 and Figure 3 illustrate the epoch-wise trends of Coverage, Ranking Loss and One Error for CCMN, L-UW, GDF, $R_u(g)$, and ComRank, for both uniform complementary labels and biased complementary labels. As observed, ComRank consistently outperforms other methods, displaying the most rapid and least pronounced decline in Coverage, One Error and Ranking Loss. Notably, GDF, $R_u(g)$, and CCMN exhibit slight instability, with fluctuations emerging either at the initial stages or towards the end of the descent. In contrast, ComRank maintains remarkable stability, underscoring its effectiveness in the gradient descent process whenever complementary labels are selected uniformly or biasedly.

H Surrogate Loss Ablation.

The Table 10 reports the *Average Precision* from an ablation study on different surrogate losses under uniform complementary labels. The surrogate losses include:

Log loss: $\tilde{\mathcal{L}}(\mathbf{g}(\mathbf{x}), \bar{y}) = \sum_{k=1, k \neq \bar{y}}^K \log(1 + (\mathbf{g}_{\bar{y}}(\mathbf{x}) - \mathbf{g}_k(\mathbf{x})))$.

Sigmoid loss: $\tilde{\mathcal{L}}(\mathbf{g}(\mathbf{x}), \bar{y}) = \sum_{k=1, k \neq \bar{y}}^K (\mathbf{g}_{\bar{y}}(\mathbf{x}) - \mathbf{g}_k(\mathbf{x}))$. Since the predicted probability \mathbf{g} in MLL is already produced through an output sigmoid function to keep \mathbf{g} in $[0,1]$, we directly use the difference between \mathbf{g} values.

Softmax loss: $\tilde{\mathcal{L}}(\mathbf{g}(\mathbf{x}), \bar{y}) = \sum_{k=1, k \neq \bar{y}}^K (\mathbf{g}_{\bar{y}}^{\text{softmax}}(\mathbf{x}) - \mathbf{g}_k^{\text{softmax}}(\mathbf{x}))$. Here, $\mathbf{g}^{\text{softmax}}$ refers to the version of \mathbf{g} where the output function is changed from sigmoid to softmax.

Our method achieves competitive performance in most cases, validating its effectiveness and Bayes consistency.

Table 10: *Average Precision* of different surrogate losses with uniform complementary labels.

Surrogate loss	Log	Sigmoid	Softmax	ComRank
scene	0.419	0.687	0.677	0.785
yeast	0.712	0.732	0.720	0.729
enron	0.547	0.591	0.567	0.627
rcv1-sl	0.263	0.475	0.280	0.605
bibtex	0.453	0.649	0.450	0.677
bookmark	0.197	0.629	0.586	0.618
nuswideBoW	0.485	0.490	0.488	0.583